

Stakeholder Consultation on the Draft Guidelines on the Classification of High-Risk AI Systems under Article 6 of the AI Act

Submitted by: Meem Arafat Manab

Affiliation: Marie Skłodowska-Curie Doctoral Fellow, Ontology Engineering Group, Universidad Politécnica de Madrid (HARNESS Doctoral Network, Horizon Europe); Research Fellow (pro bono), School of Religion, Theology, and Peace Studies, Trinity College Dublin

Stakeholder category: Academia

Country of residence: ES – Spain

Consultation period: 19 May – 23 July 2026

Sections addressed: Section IV (Article 6(2) / Annex III); Article 112(1) review

Description of activities

Marie Skłodowska-Curie Doctoral Fellow, Ontology Engineering Group, Universidad Politécnica de Madrid (HARNESS Doctoral Network, Horizon Europe). Research Fellow (pro bono), School of Religion, Theology, and Peace Studies, Trinity College Dublin. PhD research on data governance in European Data Spaces; prior research and applied work on machine learning systems, GDPR compliance, and the structural relationship between AI training and data protection rights (including arXiv:2403.05592 on the right to be forgotten and machine learning, and joint work with Dr. V. González-Castro on GDPR and linguistic data degradation, IFIP Summer School 2025). Submitting in a personal academic capacity. These comments address a structural deficiency in the Article 6(3) filter mechanism that has not, to the knowledge of this respondent, been addressed in any prior stakeholder submission: the mechanism asks the regulated party to evaluate its own risk, using tools of its own choosing, with no independent verification before market entry. The comments propose a specific institutional remedy using infrastructure the AI Act itself already establishes.

Section IV – High-risk classification according to Article 6(2) - Annex III

IV.1 Rationale and approach

Question A – Paragraph 88

Para. 88 states "there is no separate or independent assessment to determine whether the AI system poses a significant risk of harm." This makes AI the only EU high-stakes domain where the regulated party certifies its own risk. The EMA does not permit pharma to self-certify safety. ECHA does not permit manufacturers to self-classify chemical hazard. The Guidelines should explain why AI is treated differently, or recommend pre-market consultation with relevant agencies.

Question A – Paragraph 117

Para. 117 describes market surveillance as the safeguard against misuse of Article 6(3). Market surveillance to stop child labor did not exist – you had to ban it. Market surveillance is essentially letting businesses off the hook. Who is the watchmen? A post-deployment authority, under-resourced, acting on complaint. The most serious harms in regulatory history were not surveilled away. They were prohibited. The Guidelines should explain why AI risk classification is different.

Question A – Paragraph 116

Para. 116 states documentation must be available "upon the request of a market surveillance authority." Availability upon request is not verification. These are not the same thing. A filing cabinet full of self-assessments is not a safety mechanism. The Guidelines conflate procedural documentation with substantive independent review. This distinction matters for the effectiveness of the classification framework as a whole.

Question B – Additional examples

The Guidelines should include a worked example illustrating how Article 6(3) self-assessment fails in practice. A provider of an employment screening tool documents in good faith that its system performs only a preparatory task under Article 6(3)(d). It designs a deployment

workflow where human rejection of the output is structurally rare. The self-assessment is formally compliant with paras. 113-115. The system is, in substance, high-risk. The Guidelines have no mechanism to detect this before deployment.

The structural problem: the Guidelines ask the regulated party to evaluate its own risk. No comparable EU regulatory domain does this. The EMA does not ask pharmaceutical companies to certify their own drug safety. ECHA does not ask chemical manufacturers to determine their own hazard classification. Nobody asks a coal mine to set its own safety evaluation. The rationale is identical: the party with the greatest commercial interest in a favourable classification cannot be the sole author of that classification. To allow otherwise risks a situation vis-a-vis Chernobyl, where an industry operating its own safety evaluation without independent external verification produced assessments that were formally documented, procedurally compliant, and catastrophically wrong.

The AI Act already mandates benchmarks for GPAI systemic risk classification (Article 68(3)(a)(ii)) and directs the Commission to encourage benchmark development for high-risk systems (Article 15(2)). The same Act that mandates benchmarks for one classification decision declines to require them for Article 6(3). The Guidelines should address this asymmetry. Article 15(2) benchmark methodologies, once developed, must become a required evidentiary component of Article 6(3) self-assessments.

IV.2 Horizontal issues applicable to Annex III use cases

Question A – Paragraph 113

Para. 113: the filter depends on provider self-assessment. No independent testing. No external verification. The provider claims the system is not high-risk. The claim is accepted.

Para. 117 identifies market surveillance as the check. It is not a check. It is a post-deployment complaint mechanism. Article 84 Union AI Testing Support Structures should be tasked with adversarial testing – including attempts to elicit the excluded high-risk outputs – as a condition of Article 71 registration.

Question A – Sandbox

Articles 57-60 establish supervised pre-market testing – regulatory sandboxes – with exit reports under Article 57(8). The Guidelines do not connect this to Article 6(3) claims. The Act builds apparatus for supervised pre-market testing but does not require providers invoking the most consequential classification exemption to use it. Sandbox exit reports should be recommended as the evidentiary basis for Article 6(3) self-assessments where a system has undergone sandbox testing.

Question B – Additional examples and use cases

a) The Guidelines contain no example of independent technical verification of an Article 6(3) self-assessment. Every example assumes the provider's characterisation is accurate. There is no counter-example. There is no check.

The AI industry already conducts benchmark evaluation and adversarial testing – including structured jailbreak testing – before every major model release. Providers publish benchmark results as a matter of commercial routine. The COMPL-AI framework (2024) applied the AI Act's own requirements as benchmarks across 12 state-of-the-art LLMs. None were fully compliant. Independent benchmark-based classification produces different results from provider self-assessment. The methodology exists. The industry uses it already. For its own products. The Commission is declining to require for regulatory purposes what industry already does for commercial ones.

Singapore's government-developed AI Verify toolkit, open-sourced in 2023, demonstrates that government-built AI testing infrastructure is institutionally feasible. The EU has substantially greater capacity through the Article 84 Union AI Testing Support Structures. The Act already builds supervised pre-market testing infrastructure – regulatory sandboxes under Articles 57-60, with exit reports under Article 57(8). The Guidelines do not connect sandbox exit reports to Article 6(3) exemption claims. They should. The question is not feasibility. It is whether the Commission considers the Article 6(3) classification decision consequential enough to verify independently. It answered yes for pharmaceuticals. For chemicals. For medical devices. The answer here is no. The Guidelines should explain why.

b) One use case not currently captured that should be in scope:

AI systems for dynamic pricing in essential services, e.g. housing rental, consumer credit, insurance, or energy, where algorithmic pricing systematically disadvantages protected groups. Harm pathway empirically documented. Not captured by any current Annex III entry. Article 7(2)(a) and (b) criteria are met.

Questions in relation to Article 112 paragraph 1 (Evaluation and Review)

Question 1 – Use cases to add to Annex III

Yes. One use case within existing areas requires addition.

AI systems used for dynamic pricing in essential services – housing rental, consumer credit, energy supply, and essential consumer services – where pricing is determined or materially influenced by signals of individual preference, prior behaviour, or inferred willingness to pay.

The Article 7(2) criteria are met. Scale: dynamic pricing systems operate across entire populations. Targeting: the persons most affected are those whose preference signals are most legible to the system. Irreversibility: once a transaction is completed at an algorithmically inflated price, the overpayment cannot be recovered. The use case is not captured by any current Annex III entry. Point 5(b) covers creditworthiness. Point 5(c) covers insurance pricing. Neither covers dynamic pricing of goods and services in response to individual behavioural signals. The gap is structural. The harm pathway is empirically documented.

Question 2 – Existing use cases to amend

Yes. All existing use cases in Annex III require amendment to address a structural deficiency in the Article 6(3) exemption mechanism.

As currently drafted, any provider may invoke Article 6(3) on the basis of a narrative self-assessment. No independent verification is required before market entry. This undermines every use case in Annex III simultaneously. The classification framework is only as strong as the mechanism for verifying exemption claims.

The amendment required – whether by modifying the Annex III use case entries, amending Article 6(3), or both – is to require that any provider invoking Article 6(3) submit their system to independent benchmark evaluation and structured adversarial testing, including jailbreak attempts against the provider's own claimed use restrictions, before the exemption is registered under Article 71.

The priority modalities are image-based and text-based AI systems. The benchmark and adversarial testing literature for these systems is the most mature. The methodology is proven. The AI industry has been designing benchmark suites and jailbreak attack protocols for these models for years. Staff the Article 84 Union AI Testing Support Structures with engineers from the industry who already possess this expertise. The technical capacity exists. The legislative choice not to require its use is the gap this amendment should close.

Question 3 – Use cases to remove

No.

Question 4 – Expand Article 5 prohibited practices

Yes. The following practice is not adequately addressed by Article 5(1)(a) or any other existing provision.

The deployment at scale of conversational AI systems with persistent cross-session memory, without prior longitudinal evidence on the effects of such systems on human autonomy and susceptibility to manipulation.

Article 5(1)(a) targets manipulation that happens below conscious awareness in a single interaction. That is not what cross-session memory does. Cross-session memory is visible. Each conversation feels normal. But the system is building a personal model of the user – their fears, preferences, habits, emotional triggers – across months of interaction. It uses that model to influence them. Not in one moment. Gradually. The user never notices the accumulation. That is the harm. It is not subliminal. It is something worse: it is intimate. Cross-session memory is what makes this possible at scale. It is already deployed. There is no peer-reviewed evidence on what it does to human autonomy over time. Deploy first, study later is not acceptable for a capability this consequential.

Question 5 – Remove Article 5 prohibitions

Not applicable.

MEEM ARAFAT MANAB

MADRID, 2026 CE