

Comment on Notice-MVAC-2026-01: Proposed GSAR Clause 552.239-7001(j) "Unbiased AI Principles"

Submitted by: Meem Arafat Manab

Affiliation:

Marie Skłodowska-Curie Doctoral Fellow, Ontology Engineering Group, Universidad Politécnica de Madrid (HARNESS Doctoral Network, Horizon Europe); Research Associate (pro bono), School of Religion, Theology, and Peace Studies, Trinity College Dublin. PhD research on data governance in European Data Spaces; Prior work on EU AI Act implementation (ADAPT Centre, Trinity College Dublin) and on the structural limits of compliance-by-documentation in AI governance.

Docket: 2026-0331

Sequence No.: 1

Notice: Notice-MVAC-2026-01

Document: 2026-12205

Date: June 19, 2026

Summary

This comment addresses paragraph (j), "Performance, Evaluation, and Remediation," specifically the "Unbiased AI principles" in (j)(1) and the Government Evaluation Rights in (j)(2)-(3). The concern is not with the goal of truthful, non-manipulated AI outputs, which is a legitimate procurement interest. The concern is that the clause, as drafted, asks contractors to satisfy an undefined and internally unstable standard ("ideological neutrality"), then pairs that standard with a unilateral, non-transparent enforcement mechanism. The combination creates a compliance obligation that cannot be verified by the contractor in advance and cannot be contested by the contractor after the fact – which is a structural defect independent of which administration or political tendency happens to operate the clause.

Comment Entries

#	Page / Section	Exact Language in Dispute	Comment	Suggested Language	Citation / Basis
1	p. 36565, (j)(1)(i)-(ii)	"The LLM must be truthful in responding to user prompts seeking factual information or analysis. The LLM must prioritize historical accuracy, scientific inquiry, and objectivity and must acknowledge uncertainty where reliable information is incomplete or contradictory." / "The LLM must be a neutral, nonpartisan tool that does not manipulate responses in favor of ideological dogmas."	This requirement rests on three unexamined assumptions, each of which fails on its own terms. First, the premise that partisanship is itself the harm to be engineered against is never argued, only assumed. Every system that selects training data, ranks sources, or decides what counts as a reliable claim is already making value-laden choices; the honest alternative to fake neutrality is disclosed position-taking, not the elimination of position. Second, "neutral" and "nonpartisan" are not pre-political descriptors. They name a specific	Strike the neutrality/nonpartisanship requirement in (j)(1)(ii) as a standalone compliance obligation. Replace it with a disclosure requirement: contractors must document the editorial and design choices (source selection criteria, RAG corpus composition, fine-tuning objectives, system prompt constraints) that	The technical impossibility of a neutral training target is a direct extension of work on the structural limits of compliance-by-documentation in AI governance: ethics requirements that cannot be operationalized produce paperwork, not the property they claim to ensure.

#	Page / Section	Exact Language in Dispute	Comment	Suggested Language	Citation / Basis
			stance: the depoliticized, institutionally deferential, "above the fray" posture that treats the absence of a stated commitment as the absence of any commitment at all. That posture has its own content and its own history; it privileges the status quo as the default that needs no defending while treating departures from it as "ideological dogma." A clause that rewards models for <i>appearing</i> to hold no position is enshrining one specific position, and doing so under a label that denies it has done so, which is the hypocrisy at the center of this requirement. Third, and most concretely: "neutral" is not implementable as an engineering target. There is no loss function,	shape contested or policy-salient outputs, so that any position the system takes is visible and attributable rather than asserted away as "neutral."	

#	Page / Section	Exact Language in Dispute	Comment	Suggested Language	Citation / Basis
			training objective, or benchmark that operationalizes "neutrality." Every step in the pipeline (which sources are treated as reliable for retrieval, how contradictory claims are weighted, what uncertainty gets surfaced versus suppressed, etc.) requires a decision rule, and decision rules are exactly what "neutral" claims are supposed to avoid having. A model cannot be trained toward the absence of a position because the target is not a coherent data point in the manifold of possible outputs.		
2	p. 36565, (j)(2)(i), (iv)	"The Government reserves the right to conduct automated assessments of the LLM, as deployed and configured for	This creates a one-sided evaluation regime: the Government may judge a contractor's compliance with an undefined	Require that, upon a finding of non-compliance under (j)(1), the Government	The asymmetry here mirrors a known failure mode in data governance:

#	Page / Section	Exact Language in Dispute	Comment	Suggested Language	Citation / Basis
		government users, at any time using its own benchmarks." / "All benchmarks, test data, and methodologies developed or used by the Government for such assessments are considered Government Data. The Government is under no obligation to disclose or provide access to the underlying data, methodologies, or systems, with the exception of data or methodologies used as the basis for an adverse action under paragraph (j)(3)(ii)."	neutrality standard using undisclosed criteria, and the contractor has no visibility into what "neutral" meant in practice until after an adverse finding. Pairing an undefined substantive standard (Comment 1) with a non-transparent enforcement mechanism (this paragraph) means a contractor cannot design for compliance in advance. The standard is only knowable retroactively, through enforcement actions, which is the opposite of how a workable contractual obligation should function. This is a rather Kafkaesque situation, I am afraid.	disclose the specific benchmark items, prompts, or evaluation criteria that produced the finding (which (j)(3)(ii)(B) already partially provides for adverse actions) and extend the same disclosure obligation to (j)(2) assessments generally, not only those used as the basis for an adverse action.	arrangements where one party controls both the rules and the record of compliance produce documentation, not accountability, because the regulated party can never independently verify the basis for a finding against it.
3	p. 36565, (j)(1)(i)	"The Contractor must not intentionally introduce or embed partisan or ideological judgments	The word "intentionally" is doing significant and unexamined work, I am afraid. Training data	Replace "intentionally" with a standard tied to	This is the same problem the clause's own drafters

#	Page / Section	Exact Language in Dispute	Comment	Suggested Language	Citation / Basis
		into the LLM's Data Outputs through methods such as training data selection, fine-tuning, Retrieval-Augmented Generation (RAG) references, system prompts, or other configuration methods."	selection, RAG source curation, and system prompt design are not neutral default states from which "intentional" departures occur. Every one of these choices encodes some selection, and there is no neutral baseline configuration against which "intentional" embedding of ideology can be measured. The clause implies a contractor could fail to comply through ordinary, good-faith system design choices that were never intended to encode any viewpoint, simply because all such choices are selections among alternatives.	demonstrable, documented editorial selection for partisan effect (e.g., "knowingly and for the purpose of producing partisan or ideological outputs"), and clarify that ordinary data curation, safety filtering, and RAG source selection conducted for accuracy, safety, or relevance do not constitute embedding of ideological judgment merely because they involve selection among sources.	encountered with "Material Change" and "Background Data". Both terms were given operational definitions precisely because vague terms in binding procurement clauses are unworkable; "intentionally.. .ideological" should receive the same treatment.

Closing Note

The underlying goal of paragraph (j), preventing Government use of LLMs from being silently steered by a contractor's undisclosed editorial choices, is sound procurement policy. "Neutrality" is the wrong tool for that goal. Partisanship is not, on its own, a defect to be engineered out of a system; it is an unavoidable property of any system that has to choose what to retrieve, what to weigh, and what to surface as uncertain. A clause that asks contractors to certify the absence of a position asks them to certify something no model can have and no evaluator can verify; it substitutes the appearance of neutrality for the substance of accountability, and that substitution is both the hypocrisy and the unworkability of the current text. Disclosure of editorial choices, not their suppression, is the standard that can actually be audited, enforced, and defended through changes in administration.

Respectfully submitted,

MEEM ARAFAT MANAB
MADRID, 2026 CE

**MANAB MEEM
ARAFAT -
Z2171059W**

Signed by MANAB MEEM ARAFAT
- Z2171059W
Reason: I am the author of this document
Date: 2026-07-07 15:27:32 GMT2
PDFgear 2.18