

# **Comment on the Federal Trade Commission's Proposed Policy Statement Concerning the Suppression of Accuracy in Artificial Intelligence Systems**

**Submitted by:** Meem Arafat Manab

**Affiliation:** Marie Skłodowska-Curie Doctoral Fellow, Ontology Engineering Group, Universidad Politécnica de Madrid (HARNESS Doctoral Network, Horizon Europe); Research Associate (pro bono), School of Religion, Theology, and Peace Studies, Trinity College Dublin; 2026 Emerging Leaders in Ethical AI Governance Fellow, Globethics, Geneva; 2026 Public Policy New Voices Europe Fellow, Salzburg Global Seminar

**Docket:** FTC-2026-0859

**Deadline:** 31 July 2026

1. The undersigned respectfully submits this comment in response to the Federal Trade Commission's Proposed Policy Statement Concerning the Suppression of Accuracy in Artificial Intelligence Systems, published July 1, 2026, offered in a personal academic capacity. The Commission's goal of protecting consumers from AI systems that operate contrary to their reasonable expectations is legitimate. These comments raise a structural concern: the FTC statement rests on a universal factual premise that consumers have a reasonable expectation of accuracy from AI systems, a premise that the Commission's own deception framework does not support and that the AI market does not exhibit.

2. The Commission's Policy Statement on Deception, appended to *In re Cliffdale Assocs., Inc.*, 103 F.T.C. 110, 174 (1984), states that the deception test requires evaluation "from the perspective of a consumer acting reasonably in the circumstances." The Commission has consistently applied this standard as audience-specific and context-specific. See *Lavie v. Procter & Gamble Co.*, 105 Cal. App. 4th 496, 508 (2003) ("[D]eceptiveness must be viewed in light of the intended audience"); *FTC v. Cyberspace.com LLC*, 453 F.3d 1196, 1200 (9th Cir. 2006); *POM Wonderful LLC v. FTC*, 777 F.3d 478, 490 (D.C. Cir. 2015). The proposed policy statement substitutes a universal expectation of accuracy for the audience-specific inquiry its own precedent requires, without defending that substitution.

3. The AI market is not a uniform product category admitting of a single consumer expectation. Character.AI (Mountain View, CA) serves

over 20 million monthly users interacting with fictional personas; the product is designed for character consistency, not accuracy. Replika / Luka Inc. (San Francisco, CA) markets emotional companionship, not factual accuracy. Duolingo (Pittsburgh, PA) deploys AI calibrated to the learner's level, not maximal accuracy. Latitude AI / AI Dungeon (Provo, UT) produces interactive fiction where truth value is irrelevant by design. Elicit (San Francisco, CA) is built to find and present the best available evidence against a user's position. For Elicit, accuracy in the proposed policy statement's sense is not the design goal; productive intellectual challenge is. The proposed policy statement cites Character.AI in its own footnotes as an example of the AI market, then defines the universal consumer expectation as accuracy-seeking across that same market. The statement's evidentiary foundation does not support its universalizing premise.

4. The proposed policy statement's logic, applied consistently, produces a Section 5 obligation it does not acknowledge. A company that marketed its AI as fair, bias-mitigated, or designed to avoid discriminatory outcomes, and which then removes those properties in response to this enforcement guidance without disclosing that removal, has made a material misrepresentation to the consumers who chose it for those properties. That is deception under the Commission's own 1984 standard. The statement does not acknowledge that companies have already made such marketing representations, that consumers have formed expectations on their basis, and that the proposed guidance would incentivize the withdrawal of bias-mitigating properties without disclosure; precisely the undisclosed steering the statement says it is designed to prevent.

5. The proposed policy statement states Section 5 applies "even when a company engages in a deceptive act or practice in order to comply with a state law," and that an adequate disclaimer must be prominent, persistent, and not "hidden away in fine print" (Proposed Policy Statement at 8-9). The statement applies this disclosure requirement selectively. It requires prominent, persistent disclosure when outputs are adjusted to reduce disparate impact under Colorado's AI Act. It imposes no equivalent requirement when outputs are shaped by undisclosed source selection criteria, commercial partnerships, proprietary fine-tuning objectives, or any other editorial choice that equally affects what consumers receive. A consumer has no less interest in knowing that a model's outputs reflect a commercial agreement than in knowing they reflect an equity adjustment. The Commission has not explained why one category of undisclosed output shaping triggers prominent disclosure obligations while the other does not. That asymmetry is not compelled by Section 5. It is a policy choice, and the proposed policy statement does not justify it.

6. The undersigned respectfully urges the Commission to revise the proposed policy statement to reflect the audience-specific, context-specific analysis its own precedent requires; to address the symmetric Section 5 obligations its logic produces; and to explain why the disclosure requirements it proposes apply to equity-motivated output adjustments but not to the full range of undisclosed editorial choices that equally shape what consumers receive.

MEEM ARAFAT MANAB  
MADRID, 2026 CE