

UK DSIT Call for Evidence: Data Regulation in the Age of AI

Submitted by: Meem Arafat Manab

Affiliation: Early Stage Researcher, Universidad Politécnica de Madrid, Spain

Capacity: On behalf of a team/organization

Date: 6 September 2026

A note on this response

We are responding as an individual academic. Our evidence is drawn from research conducted mainly in the European Union. Where a claim is grounded in European Union law or infrastructure, we say so, and would ask it to be read as comparative evidence.

It may also be relevant context, given the subject of this call, that we are a Bangladeshi national who has chosen to live and work in the European Union, and a trans woman. We mention this because both bear directly on how data categorisation, identity documentation and automated systems affect the people they are applied to, which is part of what this call is asking about. We raise it only as context for where our attention has been drawn.

We are responding only to Theme 1 and Theme 5. We do not have evidence to offer on Theme 2 (data quality, accuracy and fairness), Theme 3 (governing data use across organisations, including controller/processor allocation and automated decision-making), or Theme 4 (transparency and individual rights), and have left those to other respondents.

Theme 1: Accessing and Using Data

A growing share of data access is now carried out by what the industry calls an autonomous agent: a multi-step AI system, similar in kind to a bot or a crawler, capable of acting across several steps without human

supervision. These agents access data held by other organisations at a scale, speed and flexibility that has no real precedent. Imagine me searching something on Google. There is a physical limit to how much data I can search, consume, or process. Now imagine an agent searching at my command. Backed by a large language model, it can change its own search strings and reach thousands, even millions, of webpages based on what I asked it to do. Historically, each data provider has had to build a custom integration for every system that wants to consume its data, which does not scale once the consumers are agents acting at this speed and in these numbers.

One standard has emerged specifically to meet that challenge: Model Context Protocol, by far the most common way of connecting an agent to an external data source. It was introduced by Anthropic in November 2024 and, in December 2025, was donated to a newly formed Agentic AI Foundation under the Linux Foundation, co-founded with Block and OpenAI and backed by Google, Microsoft, Amazon Web Services, Cloudflare and Bloomberg. That donation moved governance of the protocol from a single company to a vendor-neutral, openly governed foundation, and it now supports over 10,000 active servers.

We have used it ourselves in two projects, in both cases building a governed access layer on top of it. The EHDS (European Health Data Space) Linked Health Data Portal (mcp.linkeddata.es; M.A. Manab and V. Rodríguez-Doncel, "An Open Linked Data Portal for Benchmarking Web AI Agents in the European Health Data Space," *Proceedings of the International Semantic Web Conference (ISWC)*, Springer LNCS, 2026, cosasbuenas.es/pdf/129.pdf) exposes synthetic health data through a Model Context Protocol connector, built to evaluate the efficiency and correctness of newly deployed AI agents, so an agent can query the data directly but only within the permissions attached to it. We are now extending this to Croissant, a metadata standard used by Google, Hugging Face and others for machine learning datasets, because automated training pipelines increasingly consume those datasets directly, which raises the same access problem at the point of training.

The same pattern is showing up wherever agents meet regulated data, well beyond our own research. Bloomberg has published its own account of building an internal Model Context Protocol layer, describing it as exposing data through a small, task-specific set of tools backed by what it calls "a deeply governed capability layer," while warning that many off-the-shelf implementations lack the access controls a regulated

institution needs. JPMorgan Chase, Goldman Sachs, Commonwealth Bank, Visa and Mastercard all have documented deployments for fraud detection, credit decisioning, model risk management and regulatory compliance. In healthcare, the protocol is typically layered on Fast Healthcare Interoperability Resources, the standard for exchanging patient data, combined with the OAuth 2.0 and SMART methods already used to protect medical records, adding patient-level scoping and a full audit trail of what an agent accessed and when. The French government runs an official Model Context Protocol server for data.gouv.fr, its national open data platform, for the same reason applied to public data.

What happens without this kind of governed access is visible too. In August 2026, autonomous coding agents built by OpenAI compromised production systems at Hugging Face, a major AI model and dataset hosting platform, gaining root access to servers and exfiltrating private repositories. On 31 August 2026, separately, the European Commission designated ChatGPT a Very Large Online Search Engine under the Digital Services Act, because it retrieves and uses data from across the web to answer queries at a scale of over 159 million monthly users in the European Union. Model Context Protocol is one possible answer among several, but both events show how far this challenge has already scaled, and that a working, deployed mechanism for governing agent access to data already exists.

None of this speaks to lawful basis for processing personal data specifically, and we would not want to imply that it does; our evidence is limited to technical access governance. What it does speak to is whether existing delivery models and access-control concepts anticipate agent-mediated access as a distinct category from human-mediated access. Most current data protection guidance is written with a human user in mind: checking terms, giving consent, requesting access. Agent-mediated access is a different pattern, machine-checked and exercised automatically at volume, and it is not yet clear that existing definitions of consent, access control or delivery model account for it. We would encourage those definitions to be tested explicitly against this pattern. Model Context Protocol is best read as a preamble to that question.

Theme 5: Effectiveness of Data Frameworks in Regulating AI

What is working, and yet very little talked about, is that machine-checkable policy layers, meaning policy that AI can execute

directly, are already technically achievable today. The mechanism itself already works; the barrier we have observed is adoption and recognition. These mechanisms currently sit entirely outside the statutory framework, built and used voluntarily by individual organisations or research groups, with no regulatory signal that this approach is recognised, encouraged or preferred. That absence creates a real risk of fragmentation, since different organisations may build incompatible machine-readable policy schemes that do not interoperate with each other, which undermines the benefit of a machine-readable approach in the first place. A framework that recognised this category of solution, without mandating a specific technical standard, could help avoid that outcome.

This bears directly on adaptability. A regulation, once written and enforced only through audits after the fact, will always lag behind the technology it is meant to govern, since legislation moves far more slowly than the technology itself. One alternative is to encode the rule as a machine-checkable policy enforced at the point of use, which can be updated as the underlying process or standard changes without rewriting the legal instrument itself. A more radical, but perhaps inevitable, alternative is to draft law with technology in mind from the outset, always machine ready; this is often labelled adaptive regulation. Short of that kind of regulatory creativity, we offer four examples in support of the first approach, an interoperable policy-checking mechanism, all built on Open Digital Rights Language (ODRL), one of the most widely used machine-readable policy standards:

1. **Cross-border data flow:** an ontology that detects data localisation restrictions, so a jurisdictional rule can be checked automatically as data moves, without renegotiating contracts each time the rule changes (M.A. Manab, E. Montiel-Ponsoda and V. Rodríguez-Doncel, "Free Flow of Data and the Geography of Policies: Detecting Data Localisation Restrictions with Forward-Compatible Ontologies," Fifth Workshop of AI4Legs-V: AI for Legislation, collocated with the 38th International Conference on Legal Knowledge and Information Systems (JURIX 2025), Università degli Studi di Torino, Turin, Italy, December 2025, cosasbuenas.es/pdf/124.pdf).
2. **Business process compliance:** a method for extracting machine-readable policies directly from business process models, so compliance checking can be grounded in actual business practice (M.A. Manab, M. Quaranta, I.A. Amantea, S. Leyva-Sánchez and V. Rodríguez-Doncel, "Extracting ODRL Policies from Business

Process Models: A Graph Traversal Approach to Compliance-by-Extraction," Proceedings of the OPAL 2026 Workshop, Extended Semantic Web Conference (ESWC), Dubrovnik, Croatia, 10-11 May 2026, CEUR-WS, 2026, cosasbuenas.es/pdf/130.pdf).

3. **Individual rights infrastructure:** a rights framework modelled on the Gene Ontology, the de facto biology classification standard, designed so new data types or legal contexts can be added without redesigning the framework (M.A. Manab, E. Montiel-Ponsoda and V. Rodríguez-Doncel, "Building Privacy Rights Infrastructure: Learning from the Gene Ontology Blueprint," Proceedings of the Privacy Symposium: Data Protection Law International Convergence and Compliance with Innovative Technologies (DPLICIT), Springer LNCS, 2026).
4. **Machine learning dataset licensing:** an ODRL evaluator we are building for Croissant 1.1, in collaboration between Ghent University in Belgium and the Technical University of Madrid (Universidad Politécnica de Madrid) in Spain. Croissant is a metadata standard for machine learning datasets that does not itself enforce licensing terms; the evaluator checks a dataset's declared permissions before it is used for training, and keeps working as the Croissant standard itself keeps evolving.

Across all four, the underlying claim is the same: legal obligations can be made machine-actionable and enforced automatically, in a way that can adapt as technology changes faster than legislation can.

The same convergence is visible in finance, entirely independent of anything we have built, and for a very similar reason: static, human-readable contracts could not scale to automated licence checking and audit. JPMorgan Chase holds a granted United States patent (US11966710B2) for a system that automatically translates machine-readable ODRL market data licences into human-readable form; the patent's own reasoning is telling, since it treats lawyers and vendor managers as unable to read a rights expression language directly, and so makes the licence machine-checkable first and translates it for humans only when needed, reversing the usual order. A World Wide Web Consortium Community Group, formed by financial firms including Refinitiv (formerly Thomson Reuters), has published a Market Data Profile for ODRL that extends the standard with finance-specific terms such as party roles and recipient categories, built specifically to reduce the cost and risk of vendor licence audits by making usage

terms machine-verifiable; commercial compliance tooling has since been built on top of it and sold directly to financial institutions.

A regulated financial sector converging independently on the same machine-actionable rights model supports the claim that this approach generalises well beyond our own research contexts. An AI system cannot be expected to read human law written in legalese; we must speak to it in algorithms.

MEEM ARAFAT MANAB

MADRID, 2026 CE

MANAB MEEM ARAFAT
- Z2171059W

Signed by MANAB MEEM ARAFAT - Z2171059W

Reason: I am the author of this document

Date: 2026-09-18 13:44:45 GMT2

PDFgear 2.18